

Aim: To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud Platforms.

Theory:

Container-based microservices architectures have revolutionized how development and operations

teams test and deploy modern software. Containers allow companies to scale and deploy applications

more efficiently, but they also introduce new challenges, adding complexity by creating a whole new

infrastructure ecosystem.

Today, both large and small software companies are deploying thousands of container instances daily.

Managing this level of complexity at scale requires advanced tools. Enter Kubernetes.

Originally developed by Google, Kubernetes is an open-source container orchestration platform designed to automate the deployment, scaling, and management of containerized applications.

Kubernetes has quickly become the de facto standard for container orchestration and is the flagship

project of the Cloud Native Computing Foundation (CNCF), supported by major players like Google,

AWS, Microsoft, IBM, Intel, Cisco, and Red Hat.

Kubernetes simplifies the deployment and operation of applications in a microservice architecture by

providing an abstraction layer over a group of hosts. This allows development teams to deploy their

applications while Kubernetes takes care of key tasks, including:

- Managing resource consumption by applications or teams
- Distributing application load evenly across the infrastructure
- Automatically load balancing requests across multiple instances of an application
- Monitoring resource usage to prevent applications from exceeding resource limits and automatically restarting them if needed
- Moving application instances between hosts when resources are low or if a host fails
- Automatically utilizing additional resources when new hosts are added to the cluster
- Facilitating canary deployments and rollbacks with ease

Necessary Requirements:

- EC2 Instance: The experiment required launching a t2.medium EC2 instance with 2 CPUs, as Kubernetes demands sufficient resources for effective functioning.

- Minimum Requirements:

- Instance Type: t2.medium
- CPUs: 2
- Memory: Adequate for container orchestration.

Create 2 Security Groups for Master and Nodes and add the following rules inbound rules in those Groups.

Master:

Type	Protocol	Port range	Source	Description - optional	
HTTP	TCP	80	Anywhere-I...	0.0.0.0/0	Delete
All traffic	All	All	Anywhere-I...	0.0.0.0/0	Delete
Custom TCP	TCP	6443	Anywhere-I...	0.0.0.0/0	Delete
Custom TCP	TCP	10250	Anywhere-I...	0.0.0.0/0	Delete
Custom TCP	TCP	10251	Anywhere-I...	0.0.0.0/0	Delete
All TCP	TCP	0 - 65535	Anywhere-I...	0.0.0.0/0	Delete
Custom TCP	TCP	10252	Anywhere-I...	0.0.0.0/0	Delete
SSH	TCP	22	Anywhere-I...	0.0.0.0/0	Delete
Add rule					

Node :

Type	Protocol	Port range	Source	Description - optional	
All traffic	All	All	Anywhere-I...	0.0.0.0/0	Delete
SSH	TCP	22	Anywhere-I...	0.0.0.0/0	Delete
Custom TCP	TCP	10250	Anywhere-I...	0.0.0.0/0	Delete
All TCP	TCP	0 - 65535	Anywhere-I...	0.0.0.0/0	Delete
Custom TCP	TCP	30000 - 32767	Anywhere-I...	0.0.0.0/0	Delete
HTTP	TCP	80	Anywhere-I...	0.0.0.0/0	Delete
Add rule					

Log in to your AWS Academy/personal account and launch 3 new Ec2 Instances. Select Ubuntu as AMI and t2.medium as Instance Type and create a key of type RSA with .pem extension and move the downloaded key to the new folder. We can use 3 Different keys or 1 common

key also.

Note: A minimum of 2 CPUs are required so Please select t2.medium and do not forget to stop the instance after the experiment because it is not available in the free tier.

Launch an instance [Info](#)

Amazon EC2 allows you to create virtual machines, or instances, that run on the AWS Cloud. Quickly get started by following the simple steps below.

Name and tags [Info](#)


Name

[Add additional tags](#)


▼ Application and OS Images (Amazon Machine Image) [Info](#)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below


Quick Start




Amazon Linux




macOS




Ubuntu




Windows



Red Hat



SUSE Linux



[Browse more AMIs](#)

Including AMIs from AWS, Marketplace and the Community

Amazon Machine Image (AMI)

▼ Instance type [Info](#) | [Get advice](#)

Instance type

t2.medium

Family: t2 2 vCPU 4 GiB Memory Current generation: true
On-Demand Linux base pricing: 0.0464 USD per Hour
On-Demand RHEL base pricing: 0.0752 USD per Hour
On-Demand Windows base pricing: 0.0644 USD per Hour
On-Demand SUSE base pricing: 0.1464 USD per Hour

☒ All generations

[Compare instance types](#)

Additional costs apply for AMIs with pre-installed software

▼ Key pair (login) [Info](#)

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name - required



[Create new key pair](#)

▼ Network settings Info

Edit

Network Info

vpc-0664684f0da1dcdfe

Subnet Info

No preference (Default subnet in any availability zone)

Auto-assign public IP Info

Enable

Additional charges apply when outside of free tier allowance

Firewall (security groups) Info

A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

☐ Create security group

☒ Select existing security group

Common security groups Info

Select security groups ▼

30EXP3MASTER sg-0818cea5ce80e34f9 ✕

VPC: vpc-0664684f0da1dcdfe

Compare security group rules

Security groups that you add or remove here will be added to or removed from all your network interfaces.

All instances

Instances (3) Info

Find Instance by attribute or tag (case-sensitive)

All states

Last updated less than a minute ago

Refresh

Connect

Instance state

Actions

Launch instances

< 1 > ⚙

<input type="checkbox"/>	Name ↗	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4 ...	Elastic IP
<input type="checkbox"/>	30EXP3NODE1	i-0bd1fc6ef532a5488	Running	t2.medium	⌚ Initializing	View alarms +	us-east-1c	ec2-35-174-171-78.co...	35.174.171.78	-
<input type="checkbox"/>	30EXP3NODE2	i-0e6894aa9f5557b90	Running	t2.medium	⌚ Initializing	View alarms +	us-east-1c	ec2-3-84-162-149.com...	3.84.162.149	-
<input type="checkbox"/>	30EXP3MASTER	i-0ea2f6230b413209a	Running	t2.medium	⌚ Initializing	View alarms +	us-east-1c	ec2-44-203-170-140.co...	44.203.170.140	-

Now open the folder in the terminal 3 times for Master, Node1& Node 2 where our .pem key is stored and paste the Example command (starting with ssh -i) in the terminal.(

Connect to instance Info

Connect to your instance i-0ea2f6230b413209a (30EXP3MASTER) using any of these options

EC2 Instance Connect

Session Manager

SSH client

EC2 serial console

Instance ID
i-0ea2f6230b413209a (30EXP3MASTER)

1. Open an SSH client.
2. Locate your private key file. The key used to launch this instance is 30-EC22-MASTER-KEY.pem
3. Run this command, if necessary, to ensure your key is not publicly viewable.
chmod 400 "30-EC22-MASTER-KEY.pem"
4. Connect to your instance using its Public DNS:
ec2-44-203-170-140.compute-1.amazonaws.com

Example:
ssh -i "30-EC22-MASTER-KEY.pem" ubuntu@ec2-44-203-170-140.compute-1.amazonaws.com

Note: In most cases, the guessed username is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.

Cancel

Connect to instance Info

Connect to your instance i-0bd1fc6ef532a5488 (30EXP3NODE1) using any of these options

EC2 Instance Connect

Session Manager

SSH client

EC2 serial console

Instance ID
i-0bd1fc6ef532a5488 (30EXP3NODE1)

1. Open an SSH client.
2. Locate your private key file. The key used to launch this instance is 30-EC22-MASTER-KEY.pem
3. Run this command, if necessary, to ensure your key is not publicly viewable.
chmod 400 "30-EC22-MASTER-KEY.pem"
4. Connect to your instance using its Public DNS:
ec2-35-174-171-78.compute-1.amazonaws.com

Example:
ssh -i "30-EC22-MASTER-KEY.pem" ubuntu@ec2-35-174-171-78.compute-1.amazonaws.com

Note: In most cases, the guessed username is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.

Cancel

All 3 instances connected successfully

```
ubuntu@ip-172-31-94-184:~$ To run a command as administrator
To run a command as administrator See "man sudo_root" for details.
See "man sudo_root" for details.
ubuntu@ip-172-31-86-209:~$ |
ubuntu@ip-172-31-81-192:~$
```

Run on Master, Node 1, and Node 2 the below commands to install and setup Docker in Master, Node1, and Node2.

```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo tee
/etc/apt/trusted.gpg.d/docker.gpg > /dev/null
sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu
$(lsb_release -cs) stable"
```

```
Fetches 29.1 MB in 4s (7394 kB/s)
Reading package lists... Done
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease:
Key is stored in legacy trusted.gpg keyring (/etc/apt/trusted.gpg),
see the DEPRECATION section in apt-key(8) for details.
ubuntu@ip-172-31-86-209:~$
```

```
sudo apt-get update
sudo apt-get install -y docker-ce
```

```
Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on
this host.
ubuntu@ip-172-31-86-209:~$ |
```

```
sudo mkdir -p /etc/docker
cat <<EOF | sudo tee /etc/docker/daemon.json
{
"exec-opts": ["native.cgroupdriver=systemd"]
}
EOF
```

```
ubuntu@ip-172-31-86-209:~$ sudo mkdir -p /etc/docker
cat <<EOF | sudo tee /etc/docker/daemon.json
{
"exec-opts": ["native.cgroupdriver=systemd"]
}
EOF
{
"exec-opts": ["native.cgroupdriver=systemd"]
}
```

```
sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker
```

```
ubuntu@ip-172-31-86-209:~$ sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker
Synchronizing state of docker.service with SysV service script with
/usr/lib/systemd/systemd-sysv-install.
Executing: /usr/lib/systemd/systemd-sysv-install enable docker
```

Run the below command to install Kubernetes.

```
curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.31/deb/Release.key | sudo gpg --dearmor -o
/etc/apt/keyrings/kubernetes-apt-keyring.gpg
echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]
https://pkgs.k8s.io/core:/stable:/v1.31/deb/ /' | sudo tee /etc/apt/sources.list.d/kubernetes.list
```

```
ubuntu@ip-172-31-86-209:~$ curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.31/deb/Release.key | sudo gpg --dearmor -o /etc/apt/keyrings/kubernetes-apt-keyring.gpg
echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]
https://pkgs.k8s.io/core:/stable:/v1.31/deb/ /' | sudo tee /etc/apt/sources.list.d/kubernetes.list
deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]
https://pkgs.k8s.io/core:/stable:/v1.31/deb/ /
```

```
sudo apt-get update
```

```
sudo apt-get install -y kubelet kubeadm kubectl
```

```
sudo apt-mark hold kubelet kubeadm kubectl
```

```
ubuntu@ip-172-31-86-209:~$ sudo apt-get update
Hit:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Hit:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease
Hit:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease
Hit:4 https://download.docker.com/linux/ubuntu noble InRelease
Get:5 https://prod-cdn.packages.k8s.io/repositories/isv:/kubernetes:/core:/stable:/v1.31/deb InRelease [1186 B]
Hit:6 http://security.ubuntu.com/ubuntu noble-security InRelease
Get:7 https://prod-cdn.packages.k8s.io/repositories/isv:/kubernetes:/core:/stable:/v1.31/deb Packages [4865 B]
Fetched 6051 B in 1s (11.7 kB/s)
Reading package lists... Done
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease:
Key is stored in legacy trusted.gpg keyring (/etc/apt/trusted.gpg),
see the DEPRECATION section in apt-key(8) for details.
```

```
ubuntu@ip-172-31-86-209:~$ sudo apt-get install -y kubelet kubeadm kubectl
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  conntrack cri-tools kubernetes-cni
The following NEW packages will be installed:
  conntrack cri-tools kubeadm kubectl kubelet kubernetes-cni
0 upgraded, 6 newly installed, 0 to remove and 143 not upgraded.
Need to get 87.4 MB of archives.
After this operation, 314 MB of additional disk space will be used.
Get:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/main amd64 conntrack amd64 1:1.4.8-1ubuntu1 [37.9 kB]
Get:2 https://prod-cdn.packages.k8s.io/repositories/isv:/kubernetes:/core:/stable:/v1.31/deb cri-tools 1.31.1-1.1 [15.7 MB]
Get:3 https://prod-cdn.packages.k8s.io/repositories/isv:/kubernetes:/core:/stable:/v1.31/deb kubeadm 1.31.1-1.1 [11.4 MB]
Get:4 https://prod-cdn.packages.k8s.io/repositories/isv:/kubernetes:/core:/stable:/v1.31/deb kubectl 1.31.1-1.1 [11.2 MB]
Get:5 https://prod-cdn.packages.k8s.io/repositories/isv:/kubernetes:/core:/stable:/v1.31/deb kubernetes-cni 1.5.1-1.1 [33.9 MB]
Get:6 https://prod-cdn.packages.k8s.io/repositories/isv:/kubernetes:/core:/stable:/v1.31/deb kubelet 1.31.1-1.1 [15.2 MB]
Fetched 87.4 MB in 1s (93.8 MB/s)
Selecting previously unselected package conntrack.
```



```
ubuntu@ip-172-31-86-209:~$ sudo apt-mark hold kubelet kubeadm kubectl
kubelet set on hold.
kubeadm set on hold.
kubectl set on hold.
```

sudo systemctl enable --now kubelet

sudo apt-get install -y containerd

```
ubuntu@ip-172-31-86-209:~$ sudo systemctl enable --now kubelet
sudo apt-get install -y containerd
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following packages were automatically installed and are no longer required:
  docker-buildx-plugin docker-ce-cli docker-ce-rootless-extras
  docker-compose-plugin libltdl7 libslirp0 pigz slirp4netns
Use 'sudo apt autoremove' to remove them.
The following additional packages will be installed:
  runc
The following packages will be REMOVED:
  containerd.io docker-ce
The following NEW packages will be installed:
  containerd runc
0 upgraded, 2 newly installed, 2 to remove and 143 not upgraded.
Need to get 47.2 MB of archives.
After this operation, 53.1 MB disk space will be freed.
Get:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates/main amd64 runc amd64 1.1.12-0ubuntu3.1 [8599 kB]
Get:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates/main amd64 containerd amd64 1.7.12-0ubuntu4.1 [38.6 MB]
Fetched 47.2 MB in 1s (91.7 MB/s)
(Reading database ... 68064 files and directories currently installed.)
Removing docker-ce (5:27.3.1-1~ubuntu.24.04~noble) ...
Removing containerd.io (1.7.22-1) ...
Selecting previously unselected package runc.
(Reading database ... 68044 files and directories currently installed.)
Preparing to unpack .../runc_1.1.12-0ubuntu3.1_amd64.deb ...
Unpacking runc (1.1.12-0ubuntu3.1) ...
Selecting previously unselected package containerd.
Preparing to unpack .../containerd_1.7.12-0ubuntu4.1_amd64.deb ...
Unpacking containerd (1.7.12-0ubuntu4.1) ...
Setting up runc (1.1.12-0ubuntu3.1) ...
Setting up containerd (1.7.12-0ubuntu4.1) ...
Processing triggers for man-db (2.12.0-4build2) ...
Scanning processes...
Scanning linux images...

Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.
```

```
sudo mkdir -p /etc/containerd
```

```
sudo containerd config default | sudo tee /etc/containerd/config.toml
```

```
ubuntu@ip-172-31-86-209:~$ sudo mkdir -p /etc/containerd
sudo containerd config default | sudo tee /etc/containerd/config.toml
ml
disabled_plugins = []
imports = []
oom_score = 0
plugin_dir = ""
required_plugins = []
root = "/var/lib/containerd"
state = "/run/containerd"
temp = ""
version = 2

[cgroup]
  path = ""

[debug]
  address = ""
  format = ""
  gid = 0
  level = ""
  uid = 0

[grpc]
  address = "/run/containerd/containerd.sock"
  gid = 0
  max_recv_message_size = 16777216
  max_send_message_size = 16777216
  tcp_address = ""
  tcp_tls_ca = ""
  tcp_tls_cert = ""
  tcp_tls_key = ""
  uid = 0

[metrics]
  address = ""
  grpc_histogram = false

[plugins]

[plugins."io.containerd.gc.v1.scheduler"]
  deletion_threshold = 0
  mutation_threshold = 100
  pause_threshold = 0.02
  schedule_delay = "0s"
  startup_delay = "100ms"

[plugins."io.containerd.grpc.v1.cri"]
  cdi_spec_dirs = ["/etc/cdi", "/var/run/cdi"]
  device_ownership_from_security_context = false
```

```
sudo systemctl restart containerd
sudo systemctl enable containerd
sudo systemctl status containerd
```

```
ubuntu@ip-172-31-86-209:~$ sudo systemctl restart containerd
sudo systemctl enable containerd
sudo systemctl status containerd
● containerd.service - containerd container runtime
   Loaded: loaded (/usr/lib/systemd/system/containerd.service; e>
   Active: active (running) since Sun 2024-09-29 12:44:38 UTC; 2>
     Docs: https://containerd.io
  Main PID: 4707 (containerd)
    Tasks: 8
   Memory: 13.3M (peak: 13.6M)
      CPU: 74ms
   CGroup: /system.slice/containerd.service
           └─4707 /usr/bin/containerd

Sep 29 12:44:38 ip-172-31-86-209 containerd[4707]: time="2024-09-2>
Sep 29 12:44:38 ip-172-31-86-209 containerd[4707]: time="2024-09-2>
Sep 29 12:44:38 ip-172-31-86-209 containerd[4707]: time="2024-09-2>
Sep 29 12:44:38 ip-172-31-86-209 containerd[4707]: time="2024-09-2>
Sep 29 12:44:38 ip-172-31-86-209 containerd[4707]: time="2024-09-2>
Sep 29 12:44:38 ip-172-31-86-209 containerd[4707]: time="2024-09-2>
Sep 29 12:44:38 ip-172-31-86-209 containerd[4707]: time="2024-09-2>
Sep 29 12:44:38 ip-172-31-86-209 containerd[4707]: time="2024-09-2>
Sep 29 12:44:38 ip-172-31-86-209 systemd[1]: Started containerd.se>
Sep 29 12:44:38 ip-172-31-86-209 containerd[4707]: time="2024-09-2>
lines 1-21/21 (END)
```

sudo apt-get install -y socat

```
ubuntu@ip-172-31-86-209:~$ sudo apt-get install -y socat
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following packages were automatically installed and are no longer required:
  docker-buildx-plugin docker-ce-cli docker-ce-rootless-extras
  docker-compose-plugin libltdl7 libslirp0 pigz slirp4netns
Use 'sudo apt autoremove' to remove them.
The following NEW packages will be installed:
  socat
0 upgraded, 1 newly installed, 0 to remove and 143 not upgraded.
Need to get 374 kB of archives.
After this operation, 1649 kB of additional disk space will be used.
.
Get:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/main amd64 socat amd64 1.8.0.0-4build3 [374 kB]
Fetched 374 kB in 0s (12.2 MB/s)
Selecting previously unselected package socat.
(Reading database ... 68108 files and directories currently installed.)
Preparing to unpack .../socat_1.8.0.0-4build3_amd64.deb ...
Unpacking socat (1.8.0.0-4build3) ...
Setting up socat (1.8.0.0-4build3) ...
Processing triggers for man-db (2.12.0-4build2) ...
Scanning processes...
Scanning linux images...

Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
```

Initialize the Kubecluster .Now Perform this Command only for Master.

sudo kubeadm init --pod-network-cidr=10.244.0.0/16

```
You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
  https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.86.209:6443 --token d7w2di.1q8m9gnsjeuytcz5 \
  --discovery-token-ca-cert-hash sha256:b86914f704964e41218a9cea2562a6c20518ad07ad7157bd901e454b9296d707
```

Run this command on master and also copy and save the Join command from above.

```
mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

```
ubuntu@ip-172-31-86-209:~$ mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config  
ubuntu@ip-172-31-86-209:~$ |
```

Run the command kubectl get nodes

```
ubuntu@ip-172-31-86-209:~$ kubectl get nodes  
NAME                STATUS    ROLES    AGE   VERSION  
ip-172-31-86-209    NotReady  control-plane  2m5s  v1.31.1  
ubuntu@ip-172-31-86-209:~$ |
```

Now Run the following command on Node 1 and Node 2 to Join to master.

```
sudo kubeadm join 172.31.86.209:6443 --token d7w2di.1q8m9gnsjeuytcz5  
--discovery-token-ca-cert-hash  
sha256:b86914f704964e41218a9cea2562a6c20518ad07ad7157bd901e454b9296d707
```

Node 1

```
ubuntu@ip-172-31-94-184:~$ sudo kubeadm join 172.31.86.209:6443 --t  
oken d7w2di.1q8m9gnsjeuytcz5 --discovery-token-ca-cert-hash sha256:  
b86914f704964e41218a9cea2562a6c20518ad07ad7157bd901e454b9296d707  
[preflight] Running pre-flight checks  
[preflight] Reading configuration from the cluster...  
[preflight] FYI: You can look at this config file with 'kubectl -n  
kube-system get cm kubeadm-config -o yaml'  
[kubelet-start] Writing kubelet configuration to file "/var/lib/kub  
elet/config.yaml"  
[kubelet-start] Writing kubelet environment file with flags to file  
"/var/lib/kubelet/kubeadm-flags.env"  
[kubelet-start] Starting the kubelet  
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:1  
0248/healthz. This can take up to 4m0s  
[kubelet-check] The kubelet is healthy after 501.756799ms  
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstra  
p  
  
This node has joined the cluster:  
* Certificate signing request was sent to apiserver and a response  
was received.  
* The Kubelet was informed of the new secure connection details.  
  
Run 'kubectl get nodes' on the control-plane to see this node join  
the cluster.
```

Node 2

```
ubuntu@ip-172-31-81-192:~$ sudo kubeadm join 172.31.86.209:6443 --token d7w2di.1q8m9gnsjeuytcz5 --discovery-token-ca-cert-hash sha256:b86914f704964e41218a9cea2562a6c20518ad07ad7157bd901e454b9296d707
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:10248/healthz. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 1.500697714s
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.
```

Kubectl get nodes

```
ubuntu@ip-172-31-86-209:~$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
ip-172-31-81-192	NotReady	<none>	66s	v1.31.1
ip-172-31-86-209	NotReady	control-plane	8m25s	v1.31.1
ip-172-31-94-184	NotReady	<none>	70s	v1.31.1

Since Status is NotReady we have to add a network plugin. And also we have to give the name to the nodes.

kubectl apply -f <https://docs.projectcalico.org/manifests/calico.yaml>

```
ubuntu@ip-172-31-86-209:~$ kubectl apply -f https://docs.projectcalico.org/manifests/calico.yaml
podd disruptionbudget.policy/calico-kube-controllers created
serviceaccount/calico-kube-controllers created
serviceaccount/calico-node created
configmap/calico-config created
customresourcedefinition.apiextensions.k8s.io/bgppconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/bgpppeers.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/blockaffinities.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/caliconodestatuses.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/clusterinformations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/felixconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworksets.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/hostendpoints.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamblocks.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamconfigs.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamhandles.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ippools.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ippreservations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/kubecontrollersconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networksets.crd.projectcalico.org created
clusterrole.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrole.rbac.authorization.k8s.io/calico-node created
clusterrolebinding.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrolebinding.rbac.authorization.k8s.io/calico-node created
daemonset.apps/calico-node created
deployment.apps/calico-kube-controllers created
ubuntu@ip-172-31-86-209:~$
```


sudo systemctl status kubelet

```
ubuntu@ip-172-31-86-209:~$ sudo systemctl status kubelet
● kubelet.service - kubelet: The Kubernetes Node Agent
   Loaded: loaded (/usr/lib/systemd/system/kubelet.service; enabled; vendor preset: enabled)
   Drop-In: /usr/lib/systemd/system/kubelet.service.d
            └─10-kubeadm.conf
   Active: active (running) since Sun 2024-09-29 12:49:07 UTC; 1min 45s ago
     Docs: https://kubernetes.io/docs/
   Main PID: 5723 (kubelet)
    Tasks: 11 (limit: 4676)
   Memory: 32.9M (peak: 33.4M)
      CPU: 9.861s
   CGroup: /system.slice/kubelet.service
           └─5723 /usr/bin/kubelet --bootstrap-kubeconfig=/etc/kubeconfig --config=/etc/kubernetes/kubelet.conf --enable-admission-plugins=NodeRestriction --kubeconfig=/etc/kubernetes/kubelet.conf --logtostderr --port=10250 --v=6

Sep 29 12:58:58 ip-172-31-86-209 kubelet[5723]: I0929 12:58:58.078Z [kubeletconfig] Setting kubeletconfig for node ip-172-31-86-209: <nil>
Sep 29 12:58:58 ip-172-31-86-209 kubelet[5723]: E0929 12:58:58.078Z [kubeletconfig] Failed to set kubeletconfig for node ip-172-31-86-209: no kubeletconfig found
Sep 29 12:59:02 ip-172-31-86-209 kubelet[5723]: I0929 12:59:02.024Z [kubeletconfig] Setting kubeletconfig for node ip-172-31-86-209: <nil>
Sep 29 12:59:04 ip-172-31-86-209 kubelet[5723]: I0929 12:59:04.078Z [kubeletconfig] Setting kubeletconfig for node ip-172-31-86-209: <nil>
Sep 29 12:59:04 ip-172-31-86-209 kubelet[5723]: E0929 12:59:04.078Z [kubeletconfig] Failed to set kubeletconfig for node ip-172-31-86-209: no kubeletconfig found
Sep 29 12:59:15 ip-172-31-86-209 kubelet[5723]: I0929 12:59:15.068Z [kubeletconfig] Setting kubeletconfig for node ip-172-31-86-209: <nil>
Sep 29 12:59:15 ip-172-31-86-209 kubelet[5723]: E0929 12:59:15.068Z [kubeletconfig] Failed to set kubeletconfig for node ip-172-31-86-209: no kubeletconfig found
Sep 29 12:59:30 ip-172-31-86-209 kubelet[5723]: I0929 12:59:30.065Z [kubeletconfig] Setting kubeletconfig for node ip-172-31-86-209: <nil>
Sep 29 12:59:30 ip-172-31-86-209 kubelet[5723]: E0929 12:59:30.065Z [kubeletconfig] Failed to set kubeletconfig for node ip-172-31-86-209: no kubeletconfig found
lines 1-23/23 (END)
```

Now Run command kubectl get nodes -o wide we can see Status is ready.

```
ubuntu@ip-172-31-86-209:~$ kubectl get nodes -o wide
NAME                                STATUS    ROLES    AGE     VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE             KERNEL-VERSION   CONTAINER-RUNTIME
ip-172-31-81-192                    Ready     <none>    4m17s   v1.31.1   172.31.81.192 <none>        Ubuntu 24.04 LTS     6.8.0-1012-aws   containerd://1.7.12
ip-172-31-86-209                    Ready     control-plane 11m     v1.31.1   172.31.86.209 <none>        Ubuntu 24.04 LTS     6.8.0-1012-aws   containerd://1.7.12
ip-172-31-94-184                    Ready     <none>    4m21s   v1.31.1   172.31.94.184 <none>        Ubuntu 24.04 LTS     6.8.0-1012-aws   containerd://1.7.12
```


Now to Rename run this command

Rename to Node 1: `kubectl label node ip-172-31-94-184 kubernetes.io/role=Node1`

Rename to Node 2: `kubectl label node ip-172-31-81-192 kubernetes.io/role=Node2`

```
ubuntu@ip-172-31-86-209:~$ kubectl label node ip-172-31-94-184 kubernetes.io/role=Node1
node/ip-172-31-94-184 labeled
ubuntu@ip-172-31-86-209:~$ kubectl label node ip-172-31-81-192 kubernetes.io/role=Node2
node/ip-172-31-81-192 labeled
```

run `kubectl get nodes`

```
ubuntu@ip-172-31-86-209:~$ kubectl get nodes
NAME                                STATUS    ROLES    AGE     VERSION
ip-172-31-81-192                    Ready     Node2    7m55s   v1.31.1
ip-172-31-86-209                    Ready     control-plane 15m     v1.31.1
ip-172-31-94-184                    Ready     Node1    7m59s   v1.31.1
```

Conclusion: In this experiment, we successfully set up a Kubernetes cluster with one master and two

worker nodes on AWS EC2 instances. After installing Docker, Kubernetes tools (kubelet, kubeadm,

kubectl), and containerd on all nodes, the master node was initialized and the worker nodes were

joined to the cluster. Initially, the nodes were in the NotReady state, which was resolved by installing

the Calico network plugin. We also labeled the nodes with appropriate roles (control-plane and worker).

The cluster became fully functional with all nodes in the Ready state, demonstrating the successful

configuration and orchestration of Kubernetes.